

---

## Orefox KMS - Detailed Application Overview

### What is Orefox KMS?

Orefox KMS (Knowledge Management System) is a proof-of-concept web application designed for the mining and

exploration industry. It serves as a centralized platform for managing:

- Documents (PDFs of exploration reports, geological surveys, technical documents)
- Geospatial datasets (mining projects, exploration tenements, drillholes, prospects)
- Metadata and tagging for intelligent search and retrieval
- Geospatial intelligence for location-based queries (e.g., "find all documents within 5km of this tenement")

The application demonstrates how geological/exploration data can be stored, tagged, and accessed via a clean web

interface with spatial query capabilities.

---

## Technology Stack

### Backend

- Django 5.x - Python web framework providing ORM, migrations, authentication, and admin interface
- GeoDjango - Django's geospatial extension for handling geographic data (polygons, points)
- PostgreSQL 16 + PostGIS 3.4 - Relational database with spatial extensions for geospatial queries
- MinIO - S3-compatible object storage for file uploads (PDFs, documents)
- django-storages + boto3 - Integration between Django and S3-compatible storage

### Frontend

- HTMX - Lightweight JavaScript library for dynamic page updates without heavy frameworks
- Tailwind CSS - Utility-first CSS framework loaded via CDN (no build step required)
- Vanilla JavaScript - Minimal progressive enhancement (drag-and-drop file uploads)

## Infrastructure

- Docker + Docker Compose - Containerized development environment
- Python 3.11 - Application runtime

---

## Directory Structure Breakdown

IFB398-T07-MINDS/

|

```

├── config/          # Django project configuration
|   ├── settings.py  # Main settings (database, storage, middleware)
|   ├── urls.py      # Root URL configuration
|   ├── wsgi.py       # WSGI server entry point
|   └── asgi.py       # ASGI server entry point (async)

```

|

```

├── core/           # Main Django application
|   ├── models.py   # Database models (Organisation, Process, Document, etc.)
|   ├── views.py    # Request handlers (dashboard, upload, documents)
|   ├── urls.py     # URL routing for core app
|   ├── admin.py    # Django admin customization
|   ├── forms.py    # Form definitions (DocumentForm)
|   ├── utils.py    # Utility functions (SHA-256 file hashing)
|   ├── apps.py     # App configuration
|   └── migrations/ # Database migration files

```

```

|   └─ 0001_initial.py  # Initial schema creation
|
|   └─ templates/       # HTML templates
|       └─ base.html    # Base template with Tailwind CSS
|           └─ core/
|               └─ dashboard.html  # Main dashboard with metrics and sidebar
|               └─ upload.html     # Document upload page with drag-and-drop
|               └─ projects.html   # Projects listing page
|
|   └─ infra/           # Infrastructure configuration
|       └─ web/
|           └─ Dockerfile  # Docker image for Django app
|           └─ requirements.txt  # Python dependencies
|
|   └─ media/           # User-uploaded files (local development fallback)
|       └─ docs/        # Document storage directory
|
|   └─ docker-compose.yml  # Multi-container orchestration
|   └─ manage.py        # Django management script
|   └─ seed-db.py       # Database seeding script for test data
|   └─ pyproject.toml    # Python project metadata and dependencies
|   └─ .env             # Environment variables (secrets, config)
└─ README.md           # Project documentation

```

---

## Core Components Deep Dive

### 1. Data Models (core/models.py)

The application follows a hierarchical domain model with automatic validation:

## Organisation

- Represents a mining/exploration company
- Fields: id, name, mode (EXPLORATION or MINING), created\_at, updated\_at
- Root entity - all other models reference this

## Process (may be renamed to "Campaign")

- Represents a project or mining operation under an Organisation
- Fields: id, name, organisation, mode (PROJECT or OPERATION), commodity, geom (MultiPolygon), timestamps
- Contains geospatial boundary (MultiPolygonField with SRID 4326)
- Used for spatial queries (e.g., "which documents are within this project area?")

## Prospect, Tenement, Drillhole

- Domain-specific entities for exploration activities
- All link to both Organisation and Process
- Basic structure currently, ready for expansion

## Document

- Stores file metadata and handles uploads to MinIO
- Fields: id, title, file, organisation, process, tags (ArrayField), timestamp, doc\_type, confidentiality, checksum\_sha256, created\_by, timestamps
- Key feature: SHA-256 checksum prevents duplicate uploads
- File storage: Uses MinIO via S3 API (not filesystem)

## Model Mixins (Advanced Pattern)

ValidatedChoiceModel = ChoiceValidationMixin + AutoCleanMixin

- ChoiceValidationMixin: Automatically validates choice fields against allowed values
- AutoCleanMixin: Runs full\_clean() before every save
- Used by Organisation and Process models for robust data validation

---

## 2. Views & Business Logic (core/views.py)

### Dashboard (dashboard view)

- Main landing page after login
- Displays metrics cards: project count, document count, prospect/drillhole/tenement counts
- Shows 8 most recent documents
- Uses dynamic model loading (\_get\_model() helper) to work even if models don't exist yet

### Document Upload (upload\_doc view)

- Handles GET (display form) and POST (process upload)
- Duplicate detection workflow:
  - a. Compute SHA-256 checksum of uploaded file (via sha256\_file())
  - b. Check if any document has matching checksum
  - c. If duplicate found: reject upload with error message
  - d. If unique: save to MinIO and create database record
- Form rendered with drag-and-drop enhancement (vanilla JS in template)

### Other Views

- documents: Document library with search and pagination
- prospects, drillholes, tenements: Placeholder pages (models defined but minimal implementation)
- ai\_insights: Placeholder for future AI features

- map\_view: Placeholder for geospatial map interface
- healthcheck: JSON endpoint for container health monitoring

---

### 3. Django Admin (core/admin.py)

Customized admin interface:

```
@djadmin.register(Process)
```

```
class ProcessAdmin(GISModelAdmin): # Uses GeoDjango map widget
```

```
    list_display = ("name", "mode", "commodity")
```

- ProcessAdmin: Uses GISModelAdmin to enable interactive map editing for polygon boundaries

- DocumentAdmin: Standard admin for document metadata

- Access at <http://localhost:8000/admin>

---

### 4. URL Routing

Root URLs (config/urls.py)

```
urlpatterns = [
```

```
    path("admin/", admin.site.urls), # Django admin
```

```
    path("", include("core.urls")), # Core app URLs
```

```
]
```

Core App URLs (core/urls.py)

```
urlpatterns = [
```

```
    path("", views.dashboard, name="dashboard"),
```

```
path("upload/", views.upload_doc, name="upload"),
path("documents/", views.documents, name="documents"),
path("prospects/", views.prospects, name="prospects"),
path("drillholes/", views.drillholes, name="drillholes"),
path("tenements/", views.tenements, name="tenements"),
path("ai-insights/", views.ai_insights, name="ai_insights"),
]
```

---

## 5. Templates

### Base Template (templates/base.html)

- Minimal layout with Tailwind CSS loaded from CDN
- Defines {% block content %} for child templates
- No JavaScript build step required

### Dashboard (templates/core/dashboard.html)

- Sidebar navigation with links to all modules
- Five metric cards showing counts
- Recent documents table
- User info section with admin panel link

### Upload Page (templates/core/upload.html)

- Django form with CSRF protection
- Drag-and-drop zone (vanilla JS enhancement)
- Real-time file size display
- Recent uploads list (last 20)
- Inline JavaScript for progressive enhancement (no external dependencies)

---

## 6. File Utilities (core/utils.py)

### SHA-256 File Hashing

```
def sha256_file(django_file) -> str:
    """Compute SHA-256 checksum of uploaded file in chunks"""
    h = hashlib.sha256()
    for chunk in iter(lambda: django_file.read(8192), b''):
        h.update(chunk)
    return h.hexdigest()
```

- Reads file in 8KB chunks (memory efficient for large files)
- Used for duplicate detection before saving

---

## 7. Database Configuration (config/settings.py)

### PostgreSQL + PostGIS

```
DATABASES = {
    "default": {
        "ENGINE": "django.contrib.gis.db.backends.postgis", # GeoDjango
        "NAME": env("DB_NAME"),
        "USER": env("DB_USER"),
        "PASSWORD": env("DB_PASS"),
        "HOST": env("DB_HOST"),
        "PORT": env("DB_PORT"),
    }
}
```



## MinIO Storage

```
DEFAULT_FILE_STORAGE = "storages.backends.s3boto3.S3Boto3Storage"
```

```
AWS_S3_ENDPOINT_URL = f"http://{env('MINIO_ENDPOINT')}"
```

```
AWS_STORAGE_BUCKET_NAME = env("MINIO_BUCKET")
```

- All file uploads go to MinIO (S3-compatible object storage)
- Database only stores file metadata and references

---

## 8. Docker Infrastructure

docker-compose.yml

Orchestrates 4 services:

### 1. db (PostgreSQL + PostGIS)

- Image: postgis/postgis:16-3.4
- Healthcheck: pg\_isready
- Persistent volume: db\_data
- Port: 5432

### 2. minio (Object storage)

- Image: minio/minio:latest
- Console port: 9001
- API port: 9000
- Persistent volume: minio\_data

### 3. create-bucket (One-time setup)

- Creates MinIO bucket on startup
- Sets download permissions
- Exits after completion

### 4. web (Django application)

- Built from infra/web/Dockerfile

- Runs migrations on startup
- Port: 8000
- Volume-mounted source code (hot reload in development)

Dockerfile (infra/web/Dockerfile)

FROM python:3.11-slim

# Install GDAL, GEOS, PROJ for GeoDjango

RUN apt-get update && apt-get install -y \

gdal-bin libgdal-dev \

proj-bin libproj-dev \

libgeos-dev

# Install Python dependencies

COPY requirements.txt .

RUN pip install -r requirements.txt

# Run Django development server

CMD ["bash", "-lc", "python manage.py migrate && python manage.py runserver 0.0.0.0:8000"]

Key dependencies:

- GDAL (Geospatial Data Abstraction Library) - handles geospatial raster/vector data
- GEOS (Geometry Engine) - performs spatial operations
- PROJ (Cartographic Projections Library) - coordinate transformations

---

## 9. Environment Configuration (.env)

All configuration is externalized:

# Django

DJANGO\_DEBUG=1

SECRET\_KEY=your-long-random-string

# PostgreSQL

POSTGRES\_DB=orefox

POSTGRES\_USER=postgres

POSTGRES\_PASSWORD=postgres

DB\_HOST=db # Docker service name

DB\_PORT=5432

# MinIO

MINIO\_ROOT\_USER=minio

MINIO\_ROOT\_PASSWORD=minio12345

MINIO\_BUCKET=documents

MINIO\_ENDPOINT=minio:9000 # Docker service name

MINIO\_USE\_SSL=0

Never commit secrets to Git - .gitignore excludes .env

---

## 10. Database Seeding (seed-db.py)

Development utility for generating test data:

- Uses Faker library to generate random data
- Creates Organizations and Documents with realistic names
- Inserts directly via psycopg2 (bypasses Django ORM)
- Note: Requires local psycopg2 or run via nix-shell

---

Application Flow

## Startup Sequence

### 1. Docker Compose starts services in dependency order:

- PostgreSQL starts and runs healthcheck
- MinIO starts
- Bucket creation service runs and exits
- Django web app starts (waits for db health + bucket creation)

### 2. Django initialization:

- Loads environment variables from .env
- Connects to PostgreSQL
- Runs pending migrations (python manage.py migrate)
- Starts development server on port 8000

## Document Upload Flow

### 1. User navigates to /upload/

### 2. Selects/drops file into upload zone

### 3. JavaScript displays file name and size

### 4. User fills in metadata (title, project, doc\_type, etc.)

### 5. Form submits to upload\_doc view

### 6. Django reads uploaded file and computes SHA-256

### 7. Checks database for existing document with same checksum

### 8. If unique:

- Saves file to MinIO (via S3 API)
- Creates Document record in PostgreSQL
- Redirects to upload page with success

### 9. If duplicate:

- Re-renders form with error message
- File is NOT saved

## Geospatial Query Example

```
# Find all processes within 5km of a point
from django.contrib.gis.measure import D
from django.contrib.gis.geos import Point

point = Point(-115.8, 37.2, srid=4326) # Longitude, Latitude
nearby = Process.objects.filter(geom__distance_lte=(point, D(km=5)))
```

---

## Key Design Decisions

### 1. Container-Based Development

- Why? Eliminates "works on my machine" issues
- Benefit: No local Python/PostgreSQL/MinIO installation needed
- Tradeoff: Slightly slower startup than native development

### 2. GeoDjango + PostGIS

- Why? Mining/exploration is inherently spatial
- Benefit: Native support for spatial queries (distance, contains, intersects)
- Use case: "Find all documents related to projects within 10km of this tenement"

### 3. MinIO vs. Filesystem

- Why? S3-compatible storage is production-ready
- Benefit: Easy migration to AWS S3 or other cloud storage
- Local development: MinIO runs in Docker, no cloud account needed

### 4. HTMX + Tailwind (No Build Step)

- Why? Proof-of-concept doesn't need complex frontend
- Benefit: Fast iteration, no webpack/vite configuration
- Tradeoff: Less suitable for complex single-page apps

## 5. SHA-256 Deduplication

- Why? Prevent duplicate document uploads
- Implementation: Hash computed before save, checked against existing records
- Benefit: Saves storage space and prevents data duplication

## 6. UUID Primary Keys

- Why? Better for distributed systems and data merging
- Benefit: No ID collisions when importing from multiple sources
- Tradeoff: Slightly larger index size than integers

---

## What Each File Is Responsible For

| File               | Responsibility   |
|--------------------|--|
| manage.py          | Django management CLI entry point                                    |
| config/settings.py | Django configuration (database, storage, middleware, installed apps) |
| config/urls.py     | Root URL routing (admin + core app)                                  |
| core/models.py     | Database schema definitions (Organisation, Process, Document, etc.)  |
| core/views.py      | Request handlers and business logic (dashboard, upload, search)      |
| core/admin.py      | Django admin customization (GIS map widget for Process)              |
| core/forms.py      | Form definitions for user input (DocumentForm)                       |

|                               |  |
|-------------------------------|--|
| core/utils.py                 | Utility functions (SHA-256 hashing, commented-out metaclass experiments) |
| core/urls.py                  | URL routing for core app endpoints                                       |
| templates/base.html           | HTML base layout with Tailwind CSS                                       |
| templates/core/dashboard.html | Main dashboard UI with metrics and sidebar                               |
| templates/core/upload.html    | Document upload page with drag-and-drop                                  |
| docker-compose.yml            | Multi-container orchestration (db, minio, web)                           |
| infra/web/Dockerfile          | Django container image definition  |
| infra/web/requirements.txt    | Python dependencies (Django, GeoDjango, boto3, etc.)                     |
| seed-db.py                    | Database seeding script for test data                                    |
| .env                          | Environment variables (secrets, config)                                  |

---

## Current State & Future Directions

### What's Working

✓ Docker Compose environment with PostgreSQL + PostGIS + MinIO
 ✓ Django admin with GIS map editing
 ✓ Document upload with SHA-256 deduplication
 ✓ Basic dashboard with metrics
 ✓ Data models for exploration domain

### What's Incomplete

⚠ Prospect/Drillhole/Tenement pages are placeholders
 ⚠ No geospatial map viewer in UI (only in admin)
 ⚠ AI insights page is a placeholder
 ⚠ No user authentication/authorization (uses Django's default)
 ⚠ Document search is basic (no full-text search or advanced filters)

### Potential Enhancements

- Map Interface: Add Leaflet.js for public-facing map visualization
- Spatial Search: "Find documents near this location" feature
- AI Integration: Document summarization, entity extraction
- Advanced Search: Full-text search, tag filtering, date ranges
- File Preview: PDF viewer in browser
- Bulk Import: CSV/Shapefile import for tenements/drillholes

---

This is a well-structured proof-of-concept that demonstrates the core value proposition of a geospatial knowledge

management system for the mining industry. The architecture is clean, the code follows Django best practices, and

the container-based setup makes it easy to develop and deploy.

---

---

---